

Massively Parallel Computing

Lecture 5: Tiled Matrix Multiplication with Boundary Conditions

Acknowledgement

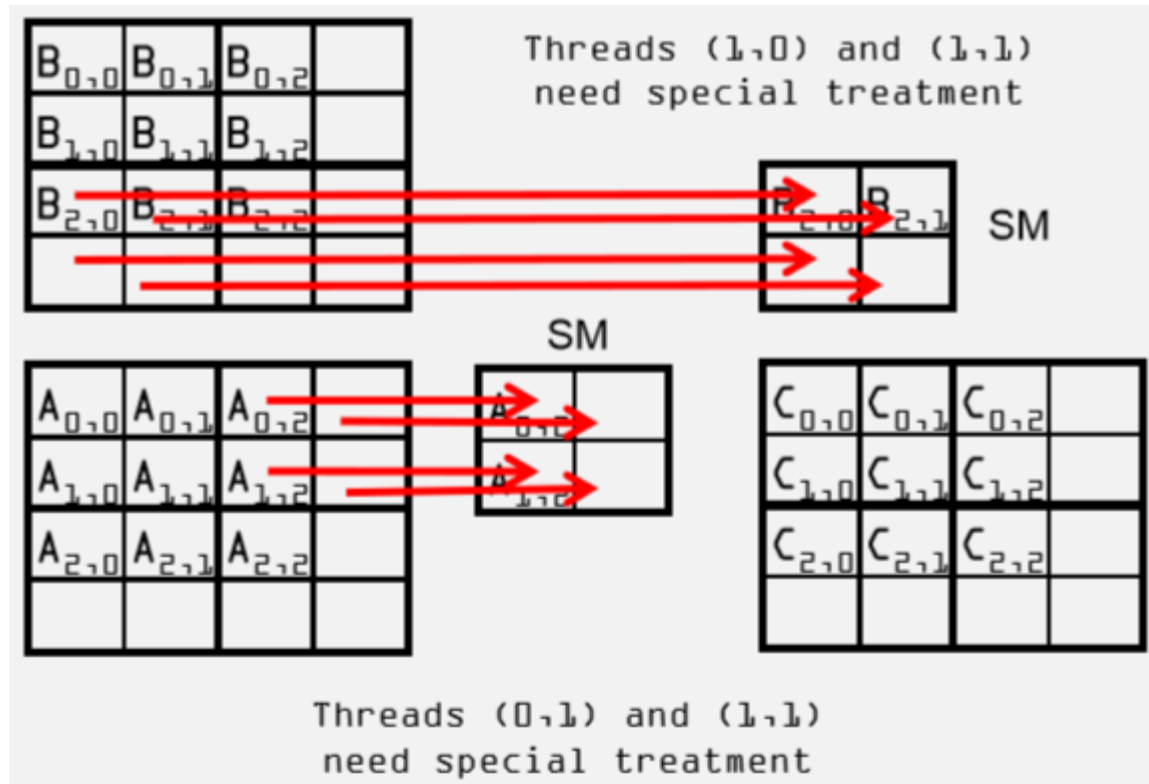
A lot of contents in this course are referred to the following sources. We deeply appreciate their effort and sharing. We promise not to use the contents for any commercial purpose.

- 1. Course materials of “Heterogeneous Parallel Programming”, University of Illinois at Urbana-Champaign, Wen-mei W. Hwu, (www.cousera.org)**
- 2. Course materials of “Massively Parallel Processors with CUDA”, Stanford University, (iTunes University)**
- 3. Course materials of “GPU Programming for High Performance Computing”, University of North Carolina at Charlotte, Barry Wilkinson**
- 4. Overheads of text book, CUDA by examples**

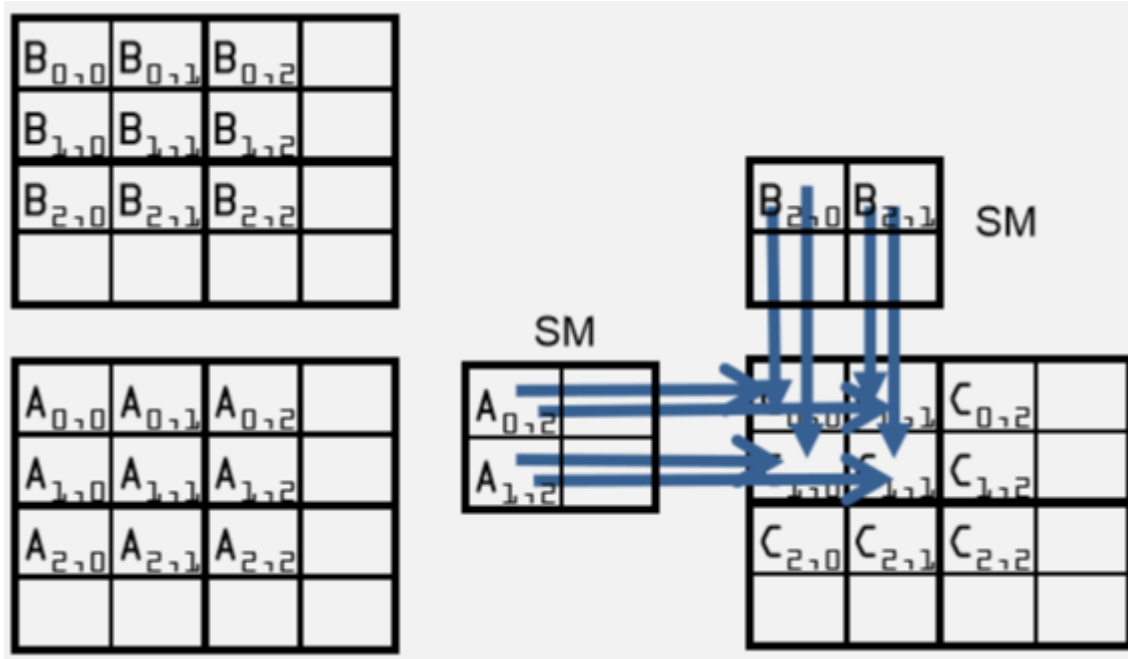
Handling Matrix of Arbitrary Size

- **The tiled matrix multiplication kernel can handle only the matrices whose dimensions are multiples of the tile width**
 - **However, real applications need to handle arbitrary sized matrices.**
 - **One could pad (add elements to) the rows and columns into multiples of the tile size, but would have significant space and data transfer time overhead.**
- **We will take a different approach**

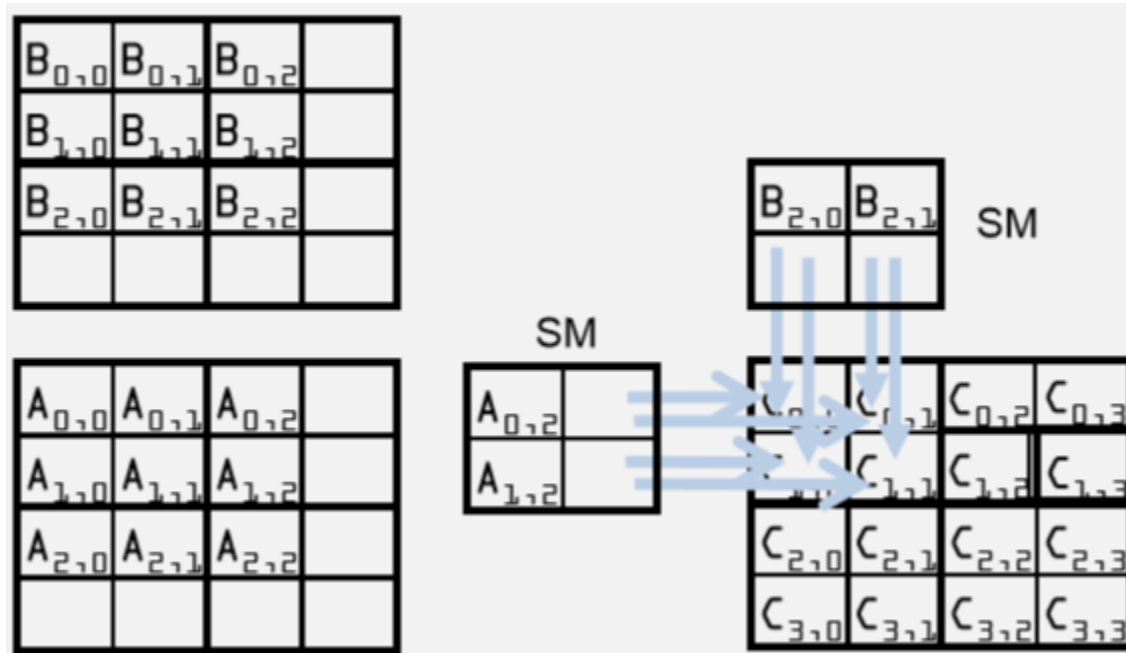
Phase 1 load for Block (0,0) for a 3x3 example



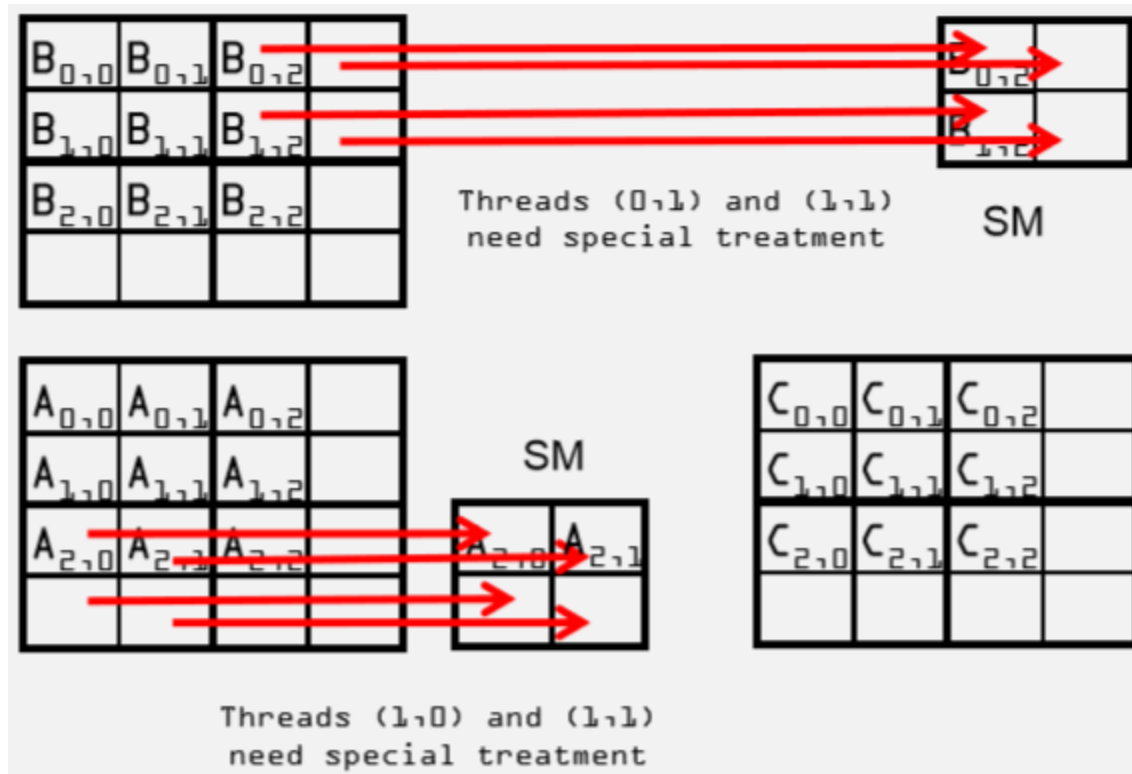
Phase 1 Use for Block (0,0) (iteration 0)



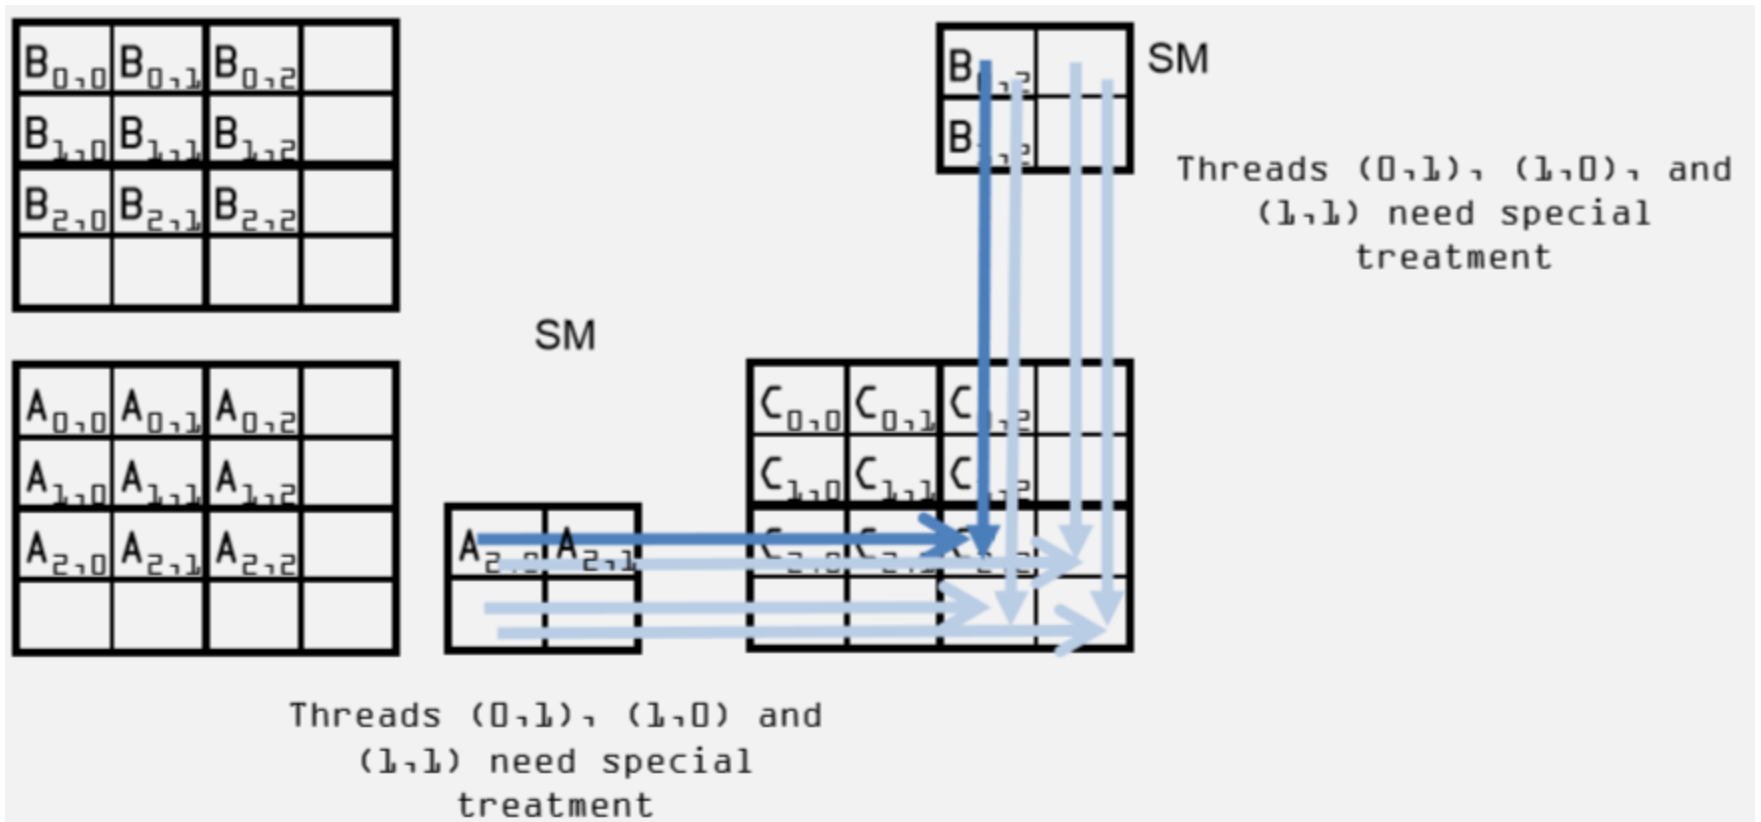
Phase 1 Use for Block (0,0) (iteration 1)



Phase 0 Load for Block (1,1) for a 3x3 example



Phase 0 use for Block(1,1) iteration 0



Phase 0 use for Block(1,1) iteration 1

$B_{0,0}$	$B_{0,1}$	$B_{0,2}$	
$B_{1,0}$	$B_{1,1}$	$B_{1,2}$	
$B_{2,0}$	$B_{2,1}$	$B_{2,2}$	

$A_{0,0}$	$A_{0,1}$	$A_{0,2}$	
$A_{1,0}$	$A_{1,1}$	$A_{1,2}$	
$A_{2,0}$	$A_{2,1}$	$A_{2,2}$	

SM

$A_{2,0}$	$A_{2,1}$

$C_{0,0}$	$C_{0,1}$	$C_{0,2}$		
$C_{1,0}$	$C_{1,1}$	$C_{1,2}$		
$C_{2,0}$	$C_{2,1}$	$C_{2,2}$		

$B_{0,2}$	
$B_{1,2}$	

SM

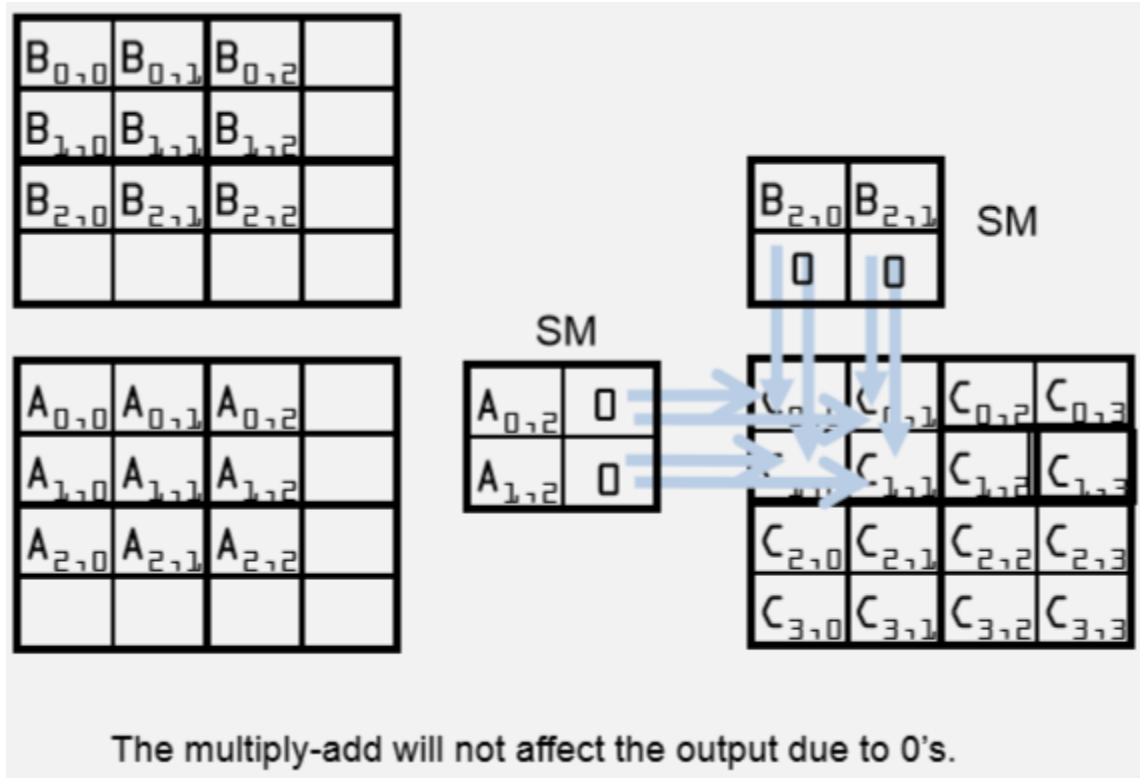
Threads (0,1), (1,0), and
(1,1) need special
treatment

Threads (0,1), (1,0) and
(1,1) need special
treatment

A “Simple” Solution


- **When a thread is to load any input element, test if it is in the valid index range**
 - **If valid, proceed to load**
 - **Else, do not load, just write a 0**
- **Rationale: a 0 value will ensure that that the multiply-add step does not affect the final value of the output element**

Phase 1 use for Block(0,0) iteration 1




Boundary Condition for Input A Tile

- Each thread loads
 - $A[\text{Row}][t * \text{TILE_WIDTH} + tx]$
 - $A[\text{Row} * \text{Width} + t * \text{TILE_WIDTH} + tx]$
- Need to test
 - $(\text{Row} < m) \ \&\& \ (t * \text{TILE_WIDTH} + tx < n)$
 - If true, load A element
 - Else , load 0



Boundary Condition for Input B Tile

- Each thread loads
 - $B[t * \text{TILE_WIDTH} + ty][\text{Col}]$
 - $B[(t * \text{TILE_WIDTH} + ty) * k + \text{Col}]$
- Need to test
 - $(t * \text{TILE_WIDTH} + ty < n) \ \&\& \ (\text{Col} < k)$
 - If true, load B element
 - Else , load 0



A “Simple” Solution (Q)

```
__global__ void MatrixMulOnDeviceWithSM(int m, int n, int k, float* A, float* B, float* C)
{
    __shared__ float ds_A[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_B[TILE_WIDTH][TILE_WIDTH];
    int bx = blockIdx.x; int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;
    int Row = by * blockDim.y + ty;
    int Col = bx * blockDim.x + tx;

    float Cvalue = 0;
    for (int t = 0; t < (n-1)/TILE_WIDTH+1; ++t) { // iterate over phases
        // load A and B tiles into shared memory
        ds_A[ty][tx] = .....
        ds_B[ty][tx] = .....
        __syncthreads();

        for (int i = 0; i < TILE_WIDTH; ++i)
            Cvalue += ds_A[ty][i] * ds_B[i][tx];
        __syncthreads();
    }
    if ( .... )
        C[Row*k+Col] = Cvalue;
}
```

Codes (Q)

```
#define LEN_M (2*1024+3)
#define LEN_N (2*1024+3)
#define LEN_K (1*1024+3)
#define TILE_WIDTH 32

int main()
{
    // Allocate and initialize the matrices A, B, C
    float * A, *B, *C, *D;
    clock_t start, end;

    A = (float*) malloc( LEN_M*LEN_N*sizeof(float) );
    B = (float*) malloc( LEN_N*LEN_K*sizeof(float) );
    C = (float*) malloc( LEN_M*LEN_K*sizeof(float) );
    D = (float*) malloc( LEN_M*LEN_K*sizeof(float) );

    for( int i=0 ; i<LEN_M*LEN_N ; i++ ) A[i] = i%3;
    for( int i=0 ; i<LEN_N*LEN_K ; i++ ) B[i] = i%4;
    for( int i=0 ; i<LEN_M*LEN_K ; i++ ) C[i] = 0.0;
    for( int i=0 ; i<LEN_M*LEN_K ; i++ ) D[i] = 0.0;

    // I/O to read the input matrices A and B
    float * dev_A, * dev_B, * dev_C;
    HANDLE_ERROR( cudaMalloc( (void**)&dev_A, LEN_M*LEN_N*sizeof(float) ));
    HANDLE_ERROR( cudaMalloc( (void**)&dev_B, LEN_N*LEN_K*sizeof(float) ));
    HANDLE_ERROR( cudaMalloc( (void**)&dev_C, LEN_M*LEN_K*sizeof(float) ));

    HANDLE_ERROR( cudaMemcpy( dev_A, A, LEN_M*LEN_N*sizeof(float)
        , cudaMemcpyHostToDevice ));
    HANDLE_ERROR( cudaMemcpy( dev_B, B, LEN_N*LEN_K*sizeof(float)
        , cudaMemcpyHostToDevice ));

    start = clock();
    // A*B on the device
    dim3 dimGrid( (LEN_K-1)/TILE_WIDTH+1, (LEN_M-1)/TILE_WIDTH+1 );
    dim3 dimBlock(TILE_WIDTH, TILE_WIDTH);
    MatrixMulOnDeviceWithSM<<<dimGrid, dimBlock>>>
        ( LEN_M, LEN_N, LEN_K, dev_A, dev_B, dev_C );

    cudaDeviceSynchronize();
```

```
end = clock();

// I/O to write the output matrix C
cudaMemcpy( C, dev_C, LEN_M*LEN_K*sizeof(long), cudaMemcpyDeviceToHost );

printf("kernel execution time: %f sec\n", (float)(end-start)/CLOCKS_PER_SEC );

MatrixMulOnHost( LEN_M, LEN_N, LEN_K, A, B, D );

printf("Check\n");
for( int i=0 ; i<LEN_M*LEN_K ; i++ ){
    if( C[i] != D[i] ) printf("Error! i=%d, C:%d, D:%d\n", i, C[i], D[i] );
}
printf("Done\n");

// Free matrices A, B, C
HANDLE_ERROR( cudaFree(dev_A) );
HANDLE_ERROR( cudaFree(dev_B) );
HANDLE_ERROR( cudaFree(dev_C) );

free(A);
free(B);
free(C);
return 0;
}

void MatrixMulOnHost(int m, int n, int k, float * A, float * B, float * C)
{
    for (int Row = 0; Row < m; ++Row){
        for (int Col = 0; Col < k; ++Col) {
            float sum = 0;
            for (int i = 0; i < n; ++i) {
                float a = A[Row * n + i];
                float b = B[Col+i*k];
                sum += a * b;
            }
            C[Row * k + Col] = sum;
        }
    }
    printf("end of matrixMulOnHost\n");
}
```